

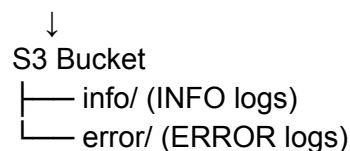
EFK Stack Deployment Documentation

Overview

This document provides a comprehensive guide to deploying the EFK (Elasticsearch, Fluentd, Kibana) stack on Kubernetes with S3 backups and log archiving.

Architecture Diagram

Application Pods → Fluentd (Log Collector) → Elasticsearch (Storage) → Kibana (Visualization)



Prerequisites

- Kubernetes cluster (k3s used in this setup)
- AWS account with S3 bucket
- AWS IAM credentials with S3 access
- kubectl configured

Step-by-Step Deployment

1. Namespace Setup

```
kubectl create namespace logging
```

2. Elasticsearch Deployment

File: `elasticsearch-deployment.yml`

```
# elasticsearch-deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: elasticsearch
  namespace: default
spec:
  replicas: 1
  selector:
    matchLabels:
      app: elasticsearch
  template:
    metadata:
      labels:
        app: elasticsearch
    spec:
      containers:
        - name: elasticsearch
          image: docker.elastic.co/elasticsearch/elasticsearch:8.11.0
          ports:
            - containerPort: 9200
            - containerPort: 9300
          env:
            - name: ES_JAVA_OPTS
              value: "-Xms512m -Xmx512m"
            - name: discovery.type
              value: single-node
            - name: xpack.security.enabled
              value: "false"
          volumeMounts:
            - name: elasticsearch-config
              mountPath: /usr/share/elasticsearch/config/elasticsearch.yml
              subPath: elasticsearch.yml
            - name: data
              mountPath: /usr/share/elasticsearch/data
      volumes:
        - name: elasticsearch-config
          configMap:
```

```
    name: elasticsearch-config
  - name: data
    emptyDir: {}
```

File: elasticsearch-service.yml

```
# elasticsearch-service.yml
apiVersion: v1
kind: Service
metadata:
  name: elasticsearch
  namespace: default
spec:
  selector:
    app: elasticsearch
  ports:
    - name: http
      port: 9200
      targetPort: 9200
    - name: transport
      port: 9300
      targetPort: 9300
  type: ClusterIP
```

Apply:

```
kubectl apply -f elasticsearch-deployment.yml
kubectl apply -f elasticsearch-service.yml
```

3. Fluentd Configuration

File: fluentd-config.yml

```
# fluentd-config.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: fluentd-config
  namespace: default
data:
  fluent.conf: |
    <system>
      log_level info
    </system>

    <source>
      @type tail
      path /var/log/containers/*.log
      pos_file /var/log/pos/fluentd-containers.log.pos
      tag kubernetes.*
      read_from_head true
      <parse>
        @type regexp
        expression /^(?<time>.+)(?<stream>stdout|stderr)[^ ]*(?<log>.*)$/
        time_format %Y-%m-%dT%H:%M:%S.%NZ
        keep_time_key true
      </parse>
    </source>

    <filter kubernetes.**>
      @type kubernetes_metadata
    </filter>

    <match kubernetes.**>
      @type elasticsearch
      host elasticsearch
      port 9200
      scheme http
      ssl_verify false
      logstash_format true
      logstash_prefix fluentd
```

```

    include_tag_key true
    type_name _doc
    flush_interval 5s

    <buffer>
      @type memory
      flush_interval 5s
      retry_forever true
    </buffer>
  </match>

  <match **>
    @type stdout
  </match>

```

Apply:

```
kubectl apply -f fluentd-config.yml
```

4. Fluentd DaemonSet

File: `fluentd-daemonset.yml`

```

# fluentd-daemonset-fixed-mounts.yml
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd
  namespace: default
  labels:
    app: fluentd
    logging: "true"
spec:
  selector:
    matchLabels:
      app: fluentd
  template:
    metadata:
      labels:

```

```
    app: fluentd
    logging: "true"
spec:
  serviceAccountName: fluentd
  tolerations:
  - key: node-role.kubernetes.io/control-plane
    operator: Exists
    effect: NoSchedule
  - key: node-role.kubernetes.io/master
    operator: Exists
    effect: NoSchedule
  containers:
  - name: fluentd
    image:
fluent/fluentd-kubernetes-daemonset:v1.16.2-debian-elasticsearch8-1.0
    env:
    - name: FLUENT_ELASTICSEARCH_HOST
      value: "elasticsearch"
    - name: FLUENT_ELASTICSEARCH_PORT
      value: "9200"
    - name: FLUENT_ELASTICSEARCH_SCHEME
      value: "http"
    - name: FLUENT_ELASTICSEARCH_SSL_VERIFY
      value: "false"
    - name: K8S_NODE_NAME
      valueFrom:
        fieldRef:
          fieldPath: spec.nodeName
    - name: FLUENT_UID
      value: "0"
  resources:
    limits:
      memory: 512Mi
    requests:
      cpu: 100m
      memory: 200Mi
  volumeMounts:
  - name: varlog
    mountPath: /var/log
  - name: varlibdockercontainers
    mountPath: /var/lib/docker/containers
    readOnly: true
  - name: dockercontainers
```

```

    mountPath: /var/log/containers
    readOnly: true
  - name: fluentd-config
    mountPath: /fluentd/etc/fluent.conf
    subPath: fluent.conf
  - name: fluentd-pos
    mountPath: /var/log/pos
terminationGracePeriodSeconds: 30
volumes:
  - name: varlog
    hostPath:
      path: /var/log
  - name: varlibdockercontainers
    hostPath:
      path: /var/lib/docker/containers
  - name: dockercontainers
    hostPath:
      path: /var/log/containers
  - name: fluentd-config
    configMap:
      name: fluentd-config
  - name: fluentd-pos
    emptyDir: {}

```

Apply:

```
kubectl apply -f fluentd-daemonset.yml
```

5. AWS S3 Credentials

```

kubectl create secret generic aws-s3-credentials \
  --namespace=logging \
  --from-literal=AWS_ACCESS_KEY_ID=YOUR_ACCESS_KEY \
  --from-literal=AWS_SECRET_ACCESS_KEY=YOUR_SECRET_KEY

```

6. Fluentd for S3

```
# fluentd-daemonset-s3.yml
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd-s3
  namespace: default
  labels:
    app: fluentd-s3
    logging: "true"
spec:
  selector:
    matchLabels:
      app: fluentd-s3
  template:
    metadata:
      labels:
        app: fluentd-s3
        logging: "true"
    spec:
      serviceAccountName: fluentd
      tolerations:
        - key: node-role.kubernetes.io/control-plane
          operator: Exists
          effect: NoSchedule
        - key: node-role.kubernetes.io/master
          operator: Exists
          effect: NoSchedule
      containers:
        - name: fluentd-s3
          image: fluent/fluentd-kubernetes-daemonset:v1.19-debian-s3-1
          env:
            - name: AWS_ACCESS_KEY_ID
              valueFrom:
                secretKeyRef:
                  name: aws-s3-credentials
                  key: AWS_ACCESS_KEY_ID
            - name: AWS_SECRET_ACCESS_KEY
              valueFrom:
                secretKeyRef:
                  name: aws-s3-credentials
                  key: AWS_SECRET_ACCESS_KEY
```



```

- name: S3_BUCKET
  value: "cloudl-efs-poc"
- name: S3_REGION
  value: "eu-north-1"
resources:
  limits:
    memory: 512Mi
  requests:
    cpu: 100m
    memory: 200Mi
volumeMounts:
- name: varlog
  mountPath: /var/log
- name: varlibdockercontainers
  mountPath: /var/lib/docker/containers
  readOnly: true
- name: fluentd-s3-config
  mountPath: /fluentd/etc/fluent.conf
  subPath: fluent.conf
- name: fluentd-pos
  mountPath: /var/log/pos
- name: fluentd-s3-buffer
  mountPath: /var/log/fluentd/s3
volumes:
- name: varlog
  hostPath:
    path: /var/log
- name: varlibdockercontainers
  hostPath:
    path: /var/lib/docker/containers
- name: fluentd-s3-config
  configMap:
    name: fluentd-s3-config
- name: fluentd-pos
  emptyDir: {}
- name: fluentd-s3-buffer
  emptyDir: {}

```

Config S3

```

# fluentd-s3-config.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: fluentd-s3-config
  namespace: default
data:
  fluent.conf: |
    <system>
      log_level info
    </system>

    <filter kubernetes.**>
      @type kubernetes_metadata
    </filter><source>
      @type tail
      path /var/log/containers/*.log
      pos_file /var/log/pos/fluentd-containers.log.pos
      tag kubernetes.*
      read_from_head true
    <parse>
      @type regexp
      expression /^(?<time>.+)(?<stream>stdout|stderr) [^ ]*(?<log>.*)$/
      time_format %Y-%m-%dT%H:%M:%S.%NZ
      keep_time_key true
    </parse>
  </source>

  # Extract log level from log content for better categorization
  <filter kubernetes.**>
    @type record_transformer
    enable_ruby true
  <record>
    log_level ${if record["log"].match(/(ERROR|Error|error)/i);
"error"; elsif
record["log"].match(/(WARN|Warn|warn|WARNING|Warning|warning)/i);
"warn"; elsif record["log"].match(/(INFO|Info|info)/i); "info"; elsif
record["log"].match(/(DEBUG|Debug|debug)/i); "debug"; else; "other";
end}
  </record>

```

```

</filter>

# Add tags to differentiate log levels for S3
<match kubernetes.**>
  @type rewrite_tag_filter
  <rule>
    key log_level
    pattern /^error$/
    tag s3.error.${tag}
  </rule>
  <rule>
    key log_level
    pattern /^warn$/
    tag s3.warn.${tag}
  </rule>
  <rule>
    key log_level
    pattern /^info$/
    tag s3.info.${tag}
  </rule>
  <rule>
    key log_level
    pattern /^debug$/
    tag s3.debug.${tag}
  </rule>
  <rule>
    key log_level
    pattern /^other$/
    tag s3.other.${tag}
  </rule>
</match>

# Send ERROR logs to S3 error folder
<match s3.error.kubernetes.**>
  @type s3
  aws_key_id "#{ENV['AWS_ACCESS_KEY_ID']}"
  aws_sec_key "#{ENV['AWS_SECRET_ACCESS_KEY']}"
  s3_bucket "cloudl-efs-poc"
  s3_region "eu-north-1"
  path "logs/error/%Y/%m/%d/"
  time_slice_format %Y%m%d%H
  <buffer time>
    @type file

```

```

    path /var/log/fluentd/s3/error
    timekey 1h
    timekey_wait 10m
    timekey_use_utc true
    chunk_limit_size 256m
</buffer>
<format>
    @type json
</format>
</match>

# Send WARN logs to S3 warn folder
<match s3.warn.kubernetes.**>
    @type s3
    aws_key_id "#{ENV['AWS_ACCESS_KEY_ID']}"
    aws_sec_key "#{ENV['AWS_SECRET_ACCESS_KEY']}"
    s3_bucket "cloudl-efs-poc"
    s3_region "eu-north-1"
    path "logs/warn/%Y/%m/%d/"
    time_slice_format %Y%m%d%H
    <buffer time>
        @type file
        path /var/log/fluentd/s3/warn
        timekey 1h
        timekey_wait 10m
        timekey_use_utc true
        chunk_limit_size 256m
    </buffer>
    <format>
        @type json
    </format>
</match>

# Send INFO logs to S3 info folder
<match s3.info.kubernetes.**>
    @type s3
    aws_key_id "#{ENV['AWS_ACCESS_KEY_ID']}"
    aws_sec_key "#{ENV['AWS_SECRET_ACCESS_KEY']}"
    s3_bucket "cloudl-efs-poc"
    s3_region "eu-north-1"
    path "logs/info/%Y/%m/%d/"
    time_slice_format %Y%m%d%H
    <buffer time>

```

```

    @type file
    path /var/log/fluentd/s3/info
    timekey 1h
    timekey_wait 10m
    timekey_use_utc true
    chunk_limit_size 256m
  </buffer>
  <format>
    @type json
  </format>
</match>

# Send DEBUG logs to S3 debug folder
<match s3.debug.kubernetes.**>
  @type s3
  aws_key_id "#{ENV['AWS_ACCESS_KEY_ID']}"
  aws_sec_key "#{ENV['AWS_SECRET_ACCESS_KEY']}"
  s3_bucket "cloudl-efs-poc"
  s3_region "eu-north-1"
  path "logs/debug/%Y/%m/%d/"
  time_slice_format %Y%m%d%H
  <buffer time>
    @type file
    path /var/log/fluentd/s3/debug
    timekey 1h
    timekey_wait 10m
    timekey_use_utc true
    chunk_limit_size 256m
  </buffer>
  <format>
    @type json
  </format>
</match>

# Send other logs to S3 other folder
<match s3.other.kubernetes.**>
  @type s3
  aws_key_id "#{ENV['AWS_ACCESS_KEY_ID']}"
  aws_sec_key "#{ENV['AWS_SECRET_ACCESS_KEY']}"
  s3_bucket "cloudl-efs-poc"
  s3_region "eu-north-1"
  path "logs/other/%Y/%m/%d/"
  time_slice_format %Y%m%d%H

```

```
<buffer time>
  @type file
  path /var/log/fluentd/s3/other
  timekey 1h
  timekey_wait 10m
  timekey_use_utc true
  chunk_limit_size 256m
</buffer>
<format>
  @type json
</format>
</match>

# Fallback for any unmatched logs
<match **>
  @type stdout
</match>
```

7. Kibana Deployment

```
# kibana-deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kibana
  namespace: default
  labels:
    app: kibana
spec:
  replicas: 1
  selector:
    matchLabels:
      app: kibana
  template:
    metadata:
      labels:
        app: kibana
    spec:
      containers:
```

```

- name: kibana
  image: docker.elastic.co/kibana/kibana:8.11.0
  ports:
    - containerPort: 5601
  env:
    - name: ELASTICSEARCH_HOSTS
      value: "http://elasticsearch:9200"
    - name: SERVER_HOST
      value: "0.0.0.0"
    - name: SERVER_NAME
      value: "kibana"
    # Since we disabled security, no credentials needed
    - name: ELASTICSEARCH_USERNAME
      value: ""
    - name: ELASTICSEARCH_PASSWORD
      value: ""
  resources:
    requests:
      cpu: 100m
      memory: 512Mi
    limits:
      memory: 1Gi

```

Kibana Service

```

# kibana-service-nodeport.yml
apiVersion: v1
kind: Service
metadata:
  name: kibana
  namespace: default
  labels:
    app: kibana
spec:
  selector:
    app: kibana
  ports:
    - port: 5601
      targetPort: 5601
      nodePort: 30601 # Custom port between 30000-32767
  type: NodePort

```

Apply:

```
kubectl apply -f kibana-deployment.yml
kubectl apply -f kibana-service.yml
```

8. RBAC Configuration

```
# fluentd-rbac.yml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: fluentd
  namespace: default
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: fluentd
  namespace: default
rules:
- apiGroups: [""]
  resources:
    - pods
    - namespaces
    - nodes
    - nodes/proxy
  verbs:
    - get
    - list
    - watch
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: fluentd
  namespace: default
roleRef:
  kind: ClusterRole
  name: fluentd
  apiGroup: rbac.authorization.k8s.io
```



```
subjects:
- kind: ServiceAccount
  name: fluentd
  namespace: default
```

Apply:

```
kubectrl apply -f fluentd-rbac.yml
```

Verification Steps

1. Check Pod Status

```
kubectrl get pods
kubectrl get svc
```

2. Test Log Generation

```
kubectrl run test-logger --image=busybox --command -- sh -c "echo 'INFO:
Test log message'; sleep 3600"
```

3. Create Index Pattern

1. Go to Stack Management → Index Patterns
2. Create pattern: `fluentd-*`
3. Time field: `@timestamp`

Backup Configuration

Elasticsearch Snapshot to S3

```
# Register S3 repository
curl -X PUT "localhost:9200/_snapshot/my_s3_backup" -H 'Content-Type: application/json' -d'
{
  "type": "s3",
  "settings": {
    "bucket": "your-backup-bucket",
    "region": "us-east-1"
  }
}'

# Create snapshot
curl -X PUT
"localhost:9200/_snapshot/my_s3_backup/snapshot_1?wait_for_completion=true"
```

Automated Backup CronJob

File: `es-backup-cronjob.yml`

```
# elasticsearch-backup-cronjob.yml
apiVersion: batch/v1
kind: CronJob
metadata:
  name: elasticsearch-backup
  namespace: default
spec:
  schedule: "0 2 * * *" # Daily at 2 AM
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: backup
              image: curlimages/curl
              env:
                - name: ES_HOST
                  value: "elasticsearch:9200"
```

```

- name: SNAPSHOT_NAME
  value: "snapshot-$(date +%Y%m%d-%H%M%S)"
command:
- /bin/sh
- -c
- |
  # Register S3 repository
  curl -X PUT "$ES_HOST/_snapshot/my_s3_repository" -H
'Content-Type: application/json' -d'
  {
    "type": "s3",
    "settings": {
      "bucket": "cloudl-efs-poc",
      "region": "eu-north-1"
    }
  }'

  # Create snapshot
  curl -X PUT
"$ES_HOST/_snapshot/my_s3_repository/$SNAPSHOT_NAME?wait_for_completion
=true"

  # Cleanup old snapshots (keep last 7)
  curl -X GET "$ES_HOST/_snapshot/my_s3_repository/_all" |
grep -o '"snapshot": "[^"]*"' | cut -d'"' -f4 | sort | head -n -7 |
xargs -I {} curl -X DELETE "$ES_HOST/_snapshot/my_s3_repository/{}"
volumeMounts:
- name: aws-credentials
  mountPath: /root/.aws
restartPolicy: OnFailure
volumes:
- name: aws-credentials
  secret:
    secretName: aws-s3-credentials

```

Create a test snapshot:

```
curl -X PUT  
"localhost:9200/_snapshot/my_s3_backup/snapshot_1?wait_for_completion=true&pretty"
```

List all snapshots in the repository:

```
curl -X GET "localhost:9200/_snapshot/my_s3_backup/_all?pretty"
```

Security Considerations

- Enable SSL/TLS for production
- Use IAM roles instead of access keys
- Enable S3 bucket encryption
- Set up network policies
- Regular security audits